

# Evaluation Service Design & Integration with Chat Agent Flow Engine

This document proposes an architecture for an automated Evaluation Service that mass-tests the **ChatAgentStatelessEngine** conversational node flow. It extends the initial sketch (shown below) with richer components, deterministic hooks, and LLM (judge) evaluation strategies.

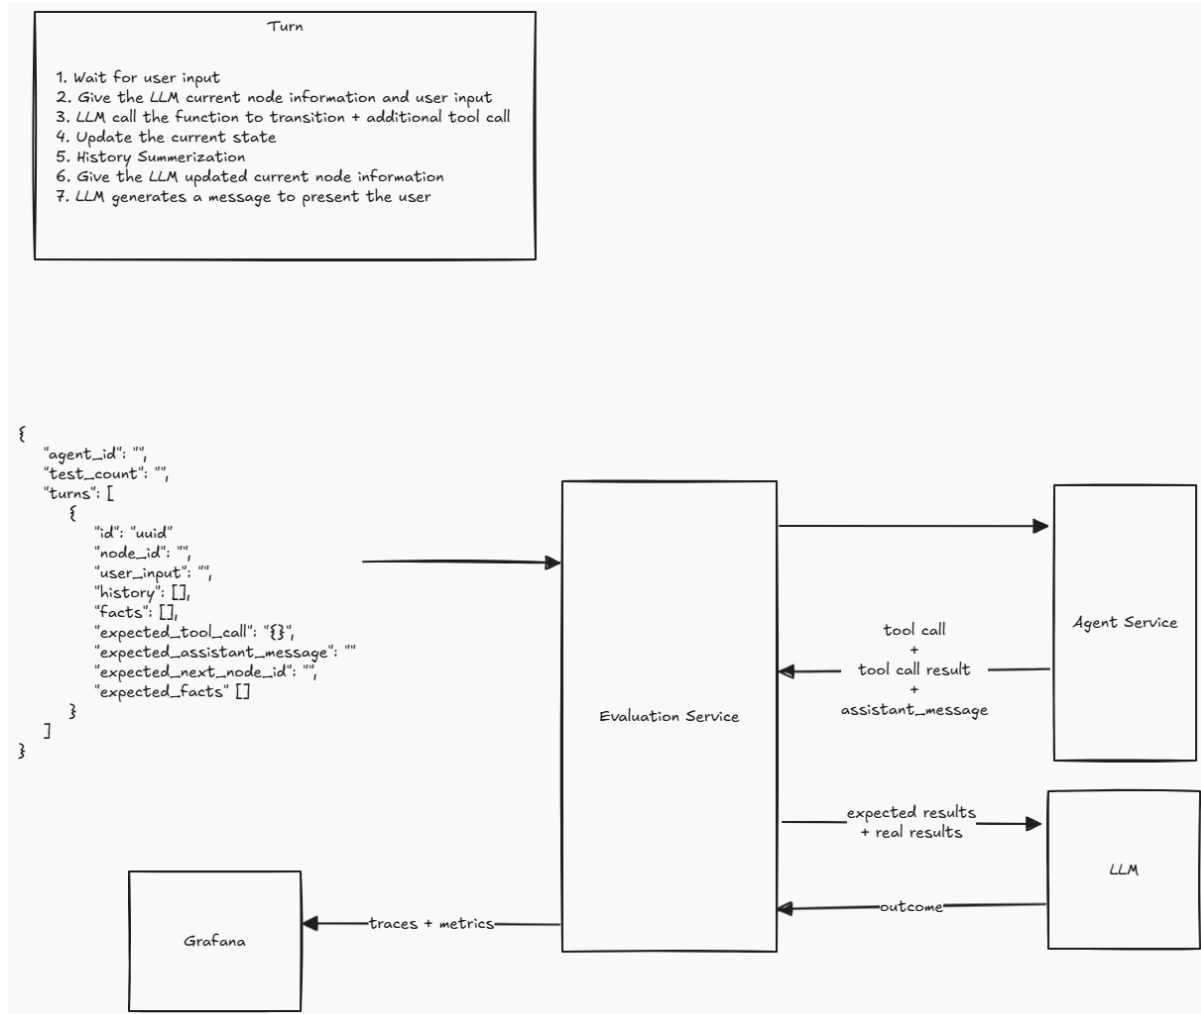


Figure 1: Initial sketch of the evaluation service architecture showing the main components and their relationships.

## Chat Agent Execution Flow

Before diving into the evaluation service, let's explore in detail the steps of the high-level turn cycle lifecycle mentioned above. A sequence diagram of the execution flow is included below to aid comprehension.

### 1. Wait for user input

- External caller supplies an **ExecutionContext** containing:
- **current\_node\_id**: ID of the flow node we are logically "in".
- **latest\_user\_message**: Raw user text for this turn.
- **history**: Dense chronological list of prior messages (role + content) kept verbatim (unsummarized source of truth).

- **memory**: Evolving **MemoryState** (turn index, extracted facts, hierarchical summary chunks, etc.).

## 2. Provide the LLM with current node information and user input

- Engine resolves / validates the current node (must not be a STOP node for execution start).
- Builds a dynamic tool schema for **process\_node** based on node type (Start vs Conversation) using parameter pydantic models.
- Generates a node descriptor (with **prompt**, node type, options, hidden "-1" stay option for conversation nodes when invoked as a tool).
- Assembles the prompt messages:
  - System root message (base instructions + enforced process\_node usage contract).
  - (Optional) Facts system message (authoritative canonical corrections / durable claims).
  - (Optional) Summary system message (hierarchical, non-verbatim aggregated prior context).
  - Recent original history window (selected with a sliding strategy to keep token budget bounded while preserving grounding).
  - Latest user message appended as a Human message.

## 3. LLM invokes the transition function (tool) and may issue additional tool calls

- Tools bound:
  - **process\_node** (always, tool\_choice pre-selected so LLM must call it first for this turn).
  - **get\_full\_history** (exposes full raw history when precise recall of earlier details is required beyond summary & window).
  - Optional global tools (e.g., **search\_knowledge\_bases**) if knowledge bases configured.
- The first assistant output is expected to be a tool call to **process\_node** with structured args:
  - Start node → requires **current\_node\_id**, **next\_node\_id**.
  - Conversation node → requires **current\_node\_id**, chosen **option\_id** (or **-1** to remain in node without advancing).
- Engine validates + executes node processing method:
  - **\_process\_start\_node**, **\_process\_conversation\_node**, (or internally **\_process\_stop\_node** if ever invoked at end).
  - Produces a **NodeProcessorResult**: **next\_node\_id**, optional **next\_node\_descriptor**, human-readable **instructions**.
- Returns a serialized tool result (descriptor + instructions) back into the LLM conversation as a **ToolMessage**.
- The LLM may then:
  - Use the descriptor/instructions to produce a final natural language response to the user.
  - Optionally call **get\_full\_history** if it needs verbatim earlier context.
  - Optionally call a knowledge base search tool for retrieval augmentation.
- A recursive / iterative loop (implemented via **process\_tool\_calls\_recursively**) consumes tool calls until the LLM emits a normal assistant message (no further tool invocation payloads).

## 4. Update the current state

- Once a final assistant message (non-tool) is produced:
  - Append the user message (this turn) to **history** (role = user).
  - Append the assistant message (role = assistant) including any debug / trace info.

- Advance `memory.turn_index` (assistant turns count) after persisting the assistant response.
- Update `current_node_id` to the resolved `next_node_id` from node processing (flow progression). If `next_node_id` is `None`, flow is considered completed.

## 5. History summarization (memory update pipeline)

- Fact extraction / merging:
  - Selects a recent window of original history (not the summarized form) and runs an extraction model prompt to pull durable factual statements or corrections.
  - Merges them into `memory.facts` (idempotent + correction aware).
- Original history summarization:
  - Checks threshold (e.g., number of turns since last chunk) to decide whether to create a new summary chunk.
  - If created, appends a compact slice summary to `memory.summary.chunks` referencing turn span metadata.
- Hierarchical re-summarization:
  - When chunk count becomes large or compression criteria met, folds existing chunks into a higher-level synthesized summary (increasing `summary.level`).
- Errors in any memory stage are caught/logged but do NOT abort the user turn (best-effort resilience).

## 6. Provide the LLM with updated current node information (next-turn preparation)

- Not executed immediately again inside the same turn; instead, the engine returns updated `ExecutionResult` containing:
  - `current_node_id` (new node for next user input).
  - Updated `history` (verbatim originals — canonical source for future summarization / retrieval).
  - Updated `memory` (facts + hierarchical summaries).
  - `flow_completed` flag (True if next node is `None` or STOP reached).
- On the *next* invocation of `execute()`, the above updated state seeds step (2) for the new user input.

## 7. LLM generates a message to present to the user

- The final assistant message is shaped by:
  - Node descriptor prompt of the *next* node (if advanced) or current node (if stayed).
  - Instructions returned by the node processor (e.g., "Reply based on prompt ...").
  - Any optional retrieval results (knowledge base tool call outputs) injected as intermediate tool messages.
  - System guardrails (facts + summary messages) ensuring consistency, brevity, and user-aligned continuity.

## Sequence Diagram

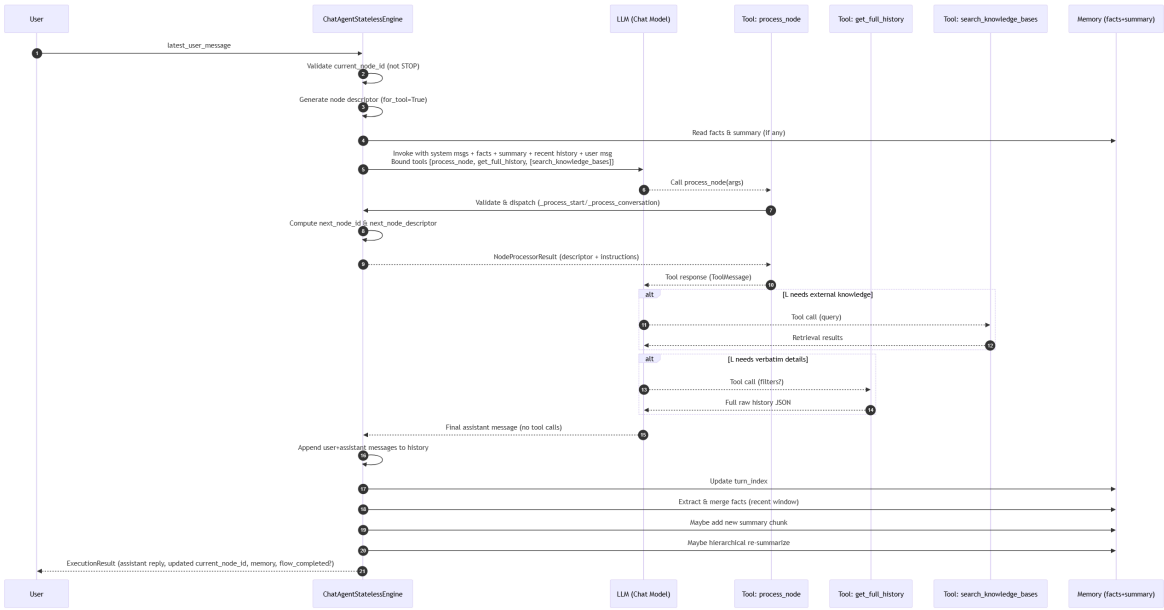


Figure 2: Sequence diagram of a single chat engine execution.

## Evaluation Service

The evaluation flow is easier to grasp visually; the sequence diagrams below illustrate the process, which the subsequent narrative explains step by step.

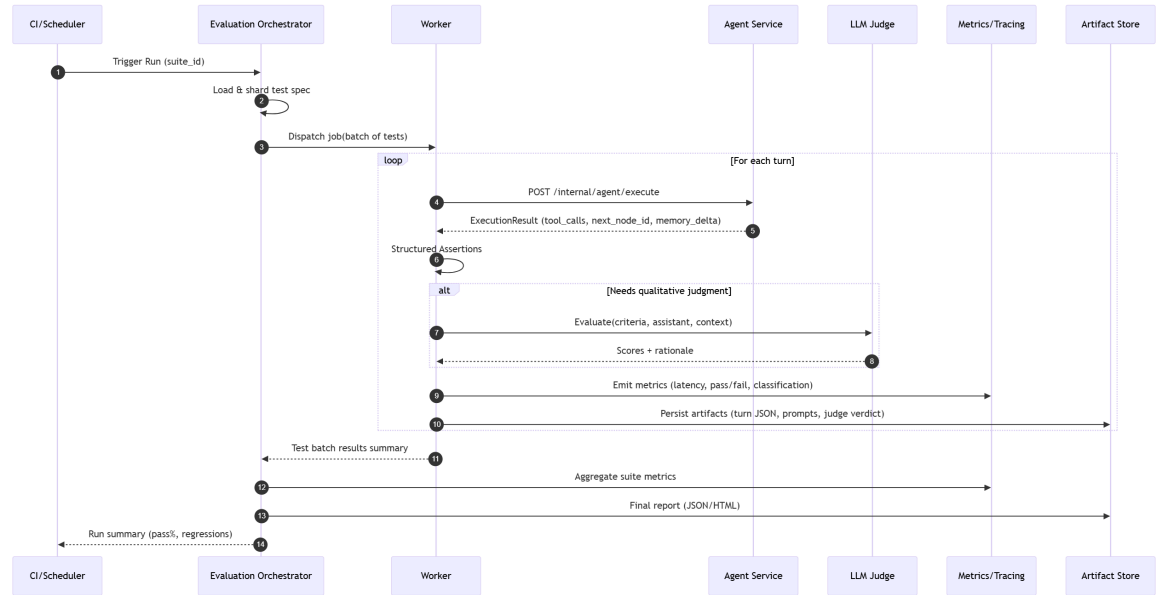


Figure 3: Sequence diagram of the evaluation engine.

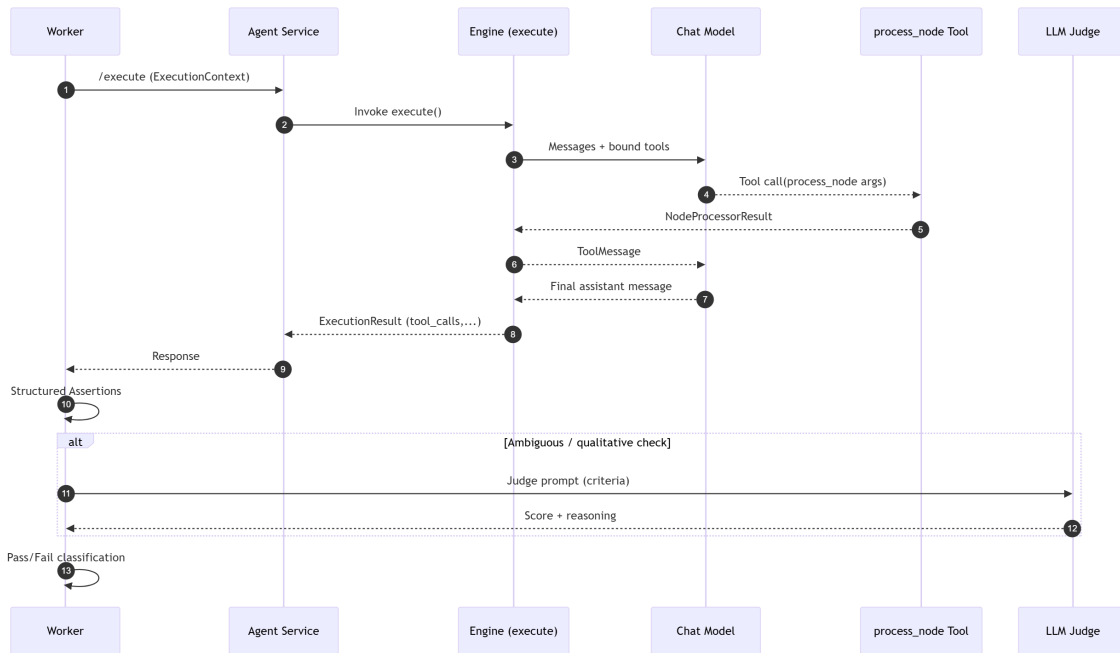


Figure 4: Sequence diagram of a single turn interaction of the evaluation engine.

## Higher-Level Components

### 1. Test Spec Repository + Scheduler (Optional)

- Test Spec Repository (Git / DB / S3) – Stores declarative JSON/YAML test suites.
- Optional scheduler to automate large-scale test execution or integrate with the CI pipeline.
- Not a near-term priority because the evaluation pipeline can be triggered manually.

### 2. Evaluation Service

- Evaluation Orchestrator – Accepts run requests; shards suites into executable jobs; manages lifecycle and aggregation.
- Execution Worker Pool – Stateless workers that execute assigned test turns against the Agent Service.
- Assertion Engine – Performs structured comparisons (tool call arguments, node transitions, fact deltas, response heuristics) before escalating to the LLM judge.
- LLM Judge Adapter – Normalizes prompts to the judge model (e.g., GPT-5 or fallback) and caches adjudications.
- Report Generator – Produces HTML / Markdown / JSON summaries (pass rates, flakiness, regression diffs).

### 3. Metrics & Tracing Layer – Emits Prometheus counters and histograms; OpenTelemetry spans correlate `test_case_id` and `conversation_id`.

### 4. Artifact Store – Persists per-turn artifacts (prompt snapshots, tool results, deltas, judge reasoning) for triage (optional).

## Required Changes to the Agent Service

The evaluation framework needs several enhancements in the Agent Service to enable deterministic, observable, and high-throughput test execution. Each required change is summarized below with its rationale and concrete implementation detail.

-- Create a new websocket endpoint (`POST /internal/agent/test`): This mirrors the public/web websocket but allows specifying model parameters (e.g., `{ "temperature": 0, "top_p": 1, "seed": <int?>, "enable_tracing": true }`). -- Execution dry-run flag: Allows running a turn without committing side effects to persistent memory. Add `dry_run: bool` to `ExecutionContext`; when true, skip persisting fact/summary updates and do not increment `turn_index`. -- Return tool call trace details: Provides intermediate visibility into reasoning for assertions. Extend `types.ExecutionResult` with `tool_calls: [{ name, args, result, latency_ms }]`, populated from the `process_tool_calls_recursively` debug flow. -- Expose full node descriptor after transition: Facilitates assertions on the next node's structure. Surface the existing tool-produced descriptor as `next_node_descriptor` in the final `ExecutionResult`. -- Fact extraction toggle: Permits disabling parts of the memory pipeline to isolate other behaviors. Add a `memory_pipeline: { enabled: bool, stages?: [facts, summarize, resummarize] }` structure in the context.

## Test Specification Schema

Proposed JSON (versioned) allowing future backward compatibility:

```
{
  "version": "v1",
  "suite_id": "flow_customer_support_v1",
  "description": "Validate onboarding flow and correction logic",
  "metadata": { "owner": "team-xyz", "priority": "P1" },
  "defaults": {
    "max_turns": 20,
    "llm_judge": { "model": "gpt-5", "criteria": ["helpfulness",
"faithfulness", "style"] }
  },
  "tests": [
    {
      "test_id": "start_handoff_path",
      "initial_node_id": "start",
      "seed": 42,
      "turns": [
        {
          "turn_id": "t1",
          "user_input": "Hello",
          "expected": {
            "tool_call": { "name": "process_node", "args.partial": {
"current_node_id": "start" } },
            "next_node_id": "conv_1",
            "assistant_contains": ["welcome"],
            "facts_add": [ { "key": "user_greeting", "value": "hello"
} ]
          }
        },
        {
          "turn_id": "t2",
          "user_input": "Option 2",
          "expected": {
            "next_node_id": "conv_2",
            "node_descriptor.options_contains": [ { "id": "opt_2" } ]
          }
        }
      ]
    }
  ]
}
```

```
    }
  }
],
"final_assertions": {
  "flow_completed": false,
  "forbidden_facts": ["incorrect_price"],
  "judge_criteria": {
    "assistant_quality_min": 0.7
  }
}
}
```

Assertion Key Types:

Key	Meaning
tool_call	Expected first tool invocation (name + partial arg match)
next_node_id	Post-turn node pointer
assistant_contains	Substrings that must appear in assistant message
assistant_not_contains	Substrings forbidden
facts_add / facts_update	Expected memory fact changes
node_descriptor.options_contains	Ensure option presence
judge_criteria	LLM-judged qualitative thresholds
flow_completed	Terminal flow expectation

## Evaluation Workflow (Turn Level)

1. Load test spec & validate schema (JSON Schema / Pydantic).
2. Initialize conversation state (empty history, optional seeded memory).
3. For each turn in `turns`:
  - Build an `ExecutionContext` with deterministic flags.
  - Invoke the Agent Service `/execute` endpoint.
  - Collect the `ExecutionResult` and trace IDs.
  - Run structured assertions:
    - Node transitions
    - Tool invocation sequence
    - Assistant text heuristics (regex / optional semantic embedding similarity)
    - Memory fact deltas
  - If any qualitative criteria are unmet or ambiguous, call the LLM judge.
  - Cache the judge verdict keyed by a hash of (`turn_input` + `assistant_output` + `criteria`) to avoid recomputation.
  - Record pass/fail classification.

- 4. After all turns: run `final_assertions`.
- 5. Emit aggregated metrics + generate artifacts.

Failure Classification:

Code	Description	Source
NODE_MISMATCH	Unexpected next node id	Structured assertion
TOOL_ARGS_MISMATCH	Tool call args deviate	Structured assertion
ASSISTANT_CONTENT	Missing/forbidden substrings	Structured assertion
FACT_DRIFT	Expected fact not added / mismatch	Memory assertion
QUALITY_JUDGE_FAIL	Judge score below threshold	LLM judge
ENGINE_ERROR	5xx / exception from Agent	Transport layer
TIMEOUT	Agent did not respond within SLA	Orchestrator
FLAKY	Non-deterministic diff across reruns	Post-processing

Retry Strategy: For transient `ENGINE_ERROR` / `TIMEOUT`, up to N (configurable, e.g., 2) retries before hard fail; judge calls are NOT retried unless network error.

Parallelization: Group tests by graph complexity and expected maximum turns to balance worker load; maintain per-conversation sequential ordering.

## 7. LLM Judge Strategy (GPT-5 or Alternatives)

Considerations when using GPT-5 (assuming availability):

- 1. Primary: `gpt-5` (when criteria complexity > threshold or prior judge failures).
- 2. Secondary: `gpt-4.1` (cost-saving) for simpler criteria (e.g., factual alignment, tone adherence).

Judge Prompt Template (Abstract):

```
SYSTEM: You are a strict evaluator. Provide JSON only.
USER: {
  "assistant_message": "...",
  "expected_criteria": {"helpfulness": true, "faithfulness": true},
  "context_excerpt": "... (node prompt / facts) ..."
}
ASSISTANT: {"scores": {"helpfulness": 0.85, "faithfulness": 0.90}, "fail_reasons":
[]}
```

Please keep in mind that the above design may change during the implementation phase.